

ENERGY LOAD PLATFORM

Technical Specification & API Reference

MVP v1.0 — For Client Review

1. Project Overview

This document outlines the technical architecture, API design, and feature scope for the Energy Load Platform MVP. It is intended for client review to confirm alignment before development proceeds.

The platform ingests raw energy consumption data in any format, normalizes it to a standard time grid, enriches it with weather and calendar context, validates it through a statistical quality pipeline, and produces an 8,760-hour annual load forecast vector with confidence intervals — the industry-standard output for energy pricing and procurement planning.

1.1 Core Value Proposition

- Upload any energy file format → receive a clean, forecast-ready annual load profile
- No manual data wrangling — AI handles format detection, unit correction, and time zone alignment
- Statistical quality assurance with automatic gap filling based on recognized usage patterns
- Output is the 8,760-hour vector (annual hours) with confidence intervals, ready for price calculation

1.2 MVP Scope Summary

Area	In MVP	Post-MVP
File Formats	CSV, Excel, MSCONS	EDI, XML, custom formats
Time Normalization	15-min UTC grid, DST handling	Configurable grid intervals
Quality Checks	Core stages (gaps, peaks, DST)	All 9 statistical stages
Weather Enrichment	API integration (read)	Historical backfill, multi-region
Calendar Enrichment	Public holidays	Custom vacation/shutdown calendars
Forecast Output	8,760 vector + confidence intervals	Multi-scenario forecasts
Frontend	Drag & Drop upload + results view	Dashboard, multi-user, roles
Auth	Single user / API key	SSO, team management

2. System Architecture

The platform is built as a containerized Python service with a lightweight frontend. All components communicate internally within a single deployable unit for the MVP, with clear module boundaries to allow future extraction into microservices.

2.1 Technology Stack

Layer	Technology	Purpose
Runtime	Python 3.11+ (Docker container)	Core processing engine
Forecasting	Python Prophet (Meta)	Time series forecasting, seasonality
Numerical	NumPy / Pandas	Vector calculation, data processing
Time Series DB	PostgreSQL + TimescaleDB	Storage and querying of time series
Object Storage	Google Cloud Storage (GCS)	Raw file storage, output artifacts
API Framework	FastAPI	REST API layer, async, auto-docs
Frontend	React (minimal)	Drag & drop upload, results display
Containerization	Docker / Docker Compose	Deployment, environment consistency
External APIs	Open-Meteo (weather), Nager.Date (holidays)	Context enrichment

2.2 Data Flow

The end-to-end data flow through the system follows these stages in sequence:

- User uploads file via drag-and-drop frontend
- File stored in GCS raw bucket, upload job created in DB
- Parser service detects format, extracts time series data
- Normalizer aligns data to 15-minute UTC grid, corrects units
- Enrichment service fetches weather and calendar data for the time range
- Quality check pipeline runs statistical tests and imputes missing values
- Forecast engine (Prophet) generates 8,760-hour annual vector
- Results stored in TimescaleDB and GCS output bucket
- Frontend polls job status and displays downloadable results

3. API Endpoints

All endpoints are prefixed with /api/v1. The API is documented automatically via FastAPI's built-in OpenAPI interface, accessible at /docs when the server is running. Authentication uses API key headers for the MVP.

3.1 Authentication

Header	Value	Required
X-API-Key	Your API key string	Yes — all endpoints
Content-Type	multipart/form-data (uploads) or application/json	Yes
Accept	application/json	Recommended

3.2 File Ingestion

POST /api/v1/upload

Accepts a raw energy data file. Detects format automatically. Returns a job ID for async tracking.

Field	Type	Description
file	multipart/form-data	The energy file (CSV, Excel, MSCONS)
timezone	string (optional)	Source timezone hint e.g. Europe/Berlin. Auto-detected if omitted.
label	string (optional)	Human-readable name for this dataset

Response (202 Accepted):

```
{ "job_id": "uuid-v4", "status": "queued", "created_at": "ISO8601" }
```

GET /api/v1/upload/{job_id}/status

Polls the processing status of an upload job.

Status Value	Meaning
queued	Job received, waiting to process
parsing	File format detection and extraction in progress
normalizing	Time alignment and unit correction running
enriching	Weather and calendar data being fetched
quality_check	Statistical validation and imputation running

Status Value	Meaning
forecasting	Prophet model generating 8,760 vector
complete	All stages done, results available
failed	Processing error — see error field in response

Response (200 OK):

```
{ "job_id": "...", "status": "forecasting", "progress_pct": 82, "stage": "forecasting", "error": null }
```

3.3 Parsing & Normalization

GET /api/v1/jobs/{job_id}/parsed

Returns the raw parsed time series after format detection and extraction, before normalization. Useful for verifying the parser understood the source file correctly.

Response Field	Type	Description
detected_format	string	e.g. CSV, MSCONS, XLSX
detected_timezone	string	Source timezone detected or provided
detected_unit	string	kW or kWh as found in source
row_count	integer	Total data rows extracted
time_range	object	start and end timestamps in source timezone
sample	array	First 10 rows of extracted data for spot-check

GET /api/v1/jobs/{job_id}/normalized

Returns the normalized time series on the 15-minute UTC grid with units corrected to kW.

Response Field	Type	Description
interval_minutes	integer	Always 15 for MVP
unit	string	Always kW after normalization
timezone	string	Always UTC after normalization
total_intervals	integer	Number of 15-min slots in the series
gap_count	integer	Gaps detected before imputation
dst_events	integer	DST transitions handled
data	array	Full normalized time series [{ts, value}]

3.4 Quality Check

GET /api/v1/jobs/{job_id}/quality-report

Returns the full statistical quality report. Each stage reports pass/fail status, findings, and actions taken.

QC Stage	What It Checks	Action if Failed
Gap Detection	Missing intervals in the time series	Imputation based on user profile
Peak/Outlier Detection	Statistically anomalous load values	Flag and optionally replace
DST Conformity	Correct handling of clock changes	Re-align affected intervals
Unit Consistency	Consistent kW vs kWh throughout file	Correct and log change
Temporal Completeness	Full year coverage check	Warn, use available data
Negative Values	Values below zero (metering errors)	Replace with imputed value
Flatline Detection	Extended identical values (sensor fault)	Flag for review
Resolution Consistency	Mixed interval resolutions	Resample to 15-min grid
Profile Plausibility	Load shape matches recognized profile	Warn if no match found

Response includes a stages array with per-stage results, an overall_pass boolean, and an imputation_count showing how many intervals were auto-filled.

3.5 Forecast Output

GET /api/v1/jobs/{job_id}/forecast

Returns the 8,760-hour annual forecast vector. This is the primary output of the platform.

Response Field	Type	Description
vector_length	integer	Always 8,760 (hours in a year)
base_year	integer	The year the forecast represents
confidence_level	float	e.g. 0.95 for 95% confidence interval
peak_load_kw	float	Maximum predicted load value
base_load_kw	float	Minimum predicted load value
annual_consumption_kwh	float	Total integrated consumption
data	array	8,760 entries: [{hour, yhat, yhat_lower, yhat_upper}]

GET /api/v1/jobs/{job_id}/forecast/download

Returns the forecast as a downloadable file.

Query Param	Options	Default
format	csv, xlsx, json	csv
include_confidence	true, false	true
timezone	Any IANA timezone string	UTC

3.6 Enrichment Data

GET /api/v1/jobs/{job_id}/enrichment

Returns the weather and calendar data that was attached to this dataset for context.

Field	Source	Description
weather.provider	Open-Meteo API	Name of weather data provider
weather.variables	Open-Meteo API	e.g. temperature_2m, solar_radiation
weather.coverage_pct	Open-Meteo API	Percentage of time range covered
calendar.country_code	Nager.Date API	e.g. DE, GB, FR
calendar.holidays	Nager.Date API	List of public holidays in the time range
calendar.holiday_count	Nager.Date API	Number of holidays detected

3.7 Health & System

Endpoint	Method	Description
GET /health	GET	Liveness check. Returns 200 OK if service is up.
GET /api/v1/status	GET	Returns service version, DB connection status, storage status.
GET /docs	GET	Auto-generated OpenAPI interactive docs (FastAPI built-in).
GET /redoc	GET	Alternative API docs in ReDoc format.

4. Key Data Models

The following describes the core data structures used throughout the platform.

4.1 Job

A Job represents a single file upload and its entire processing lifecycle.

Field	Type	Description
id	UUID	Unique job identifier
status	enum	queued parsing normalizing enriching quality_check forecasting complete failed
file_name	string	Original uploaded filename
file_size_bytes	integer	File size
gcs_raw_path	string	Path to raw file in GCS
gcs_output_path	string	Path to forecast output in GCS
created_at	timestamp	Job creation time (UTC)
completed_at	timestamp	Completion time, null until done
error_message	string	Error detail if status is failed

4.2 TimeSeries Record (TimescaleDB)

Each normalized data point is stored as a hypertable row, partitioned by time for efficient querying.

Column	Type	Description
ts	timestamptz	Timestamp (UTC, PRIMARY KEY with job_id)
job_id	UUID	Foreign key to job
value_kw	float8	Load value in kW
is_imputed	boolean	True if value was auto-filled
imputation_method	string	e.g. profile_mean, linear_interpolation

4.3 Forecast Record (TimescaleDB)

Column	Type	Description
hour_ts	timestamptz	Hourly timestamp for the forecast (UTC)
job_id	UUID	Foreign key to job
yhat	float8	Predicted load in kW
yhat_lower	float8	Lower confidence bound

Column	Type	Description
yhat_upper	float8	Upper confidence bound

5. External API Dependencies

5.1 Weather Data — Open-Meteo

Attribute	Detail
Provider	Open-Meteo (open-meteo.com)
Cost	Free tier available; no API key required for basic use
Key Variables	temperature_2m, solar_radiation, wind_speed_10m, precipitation
Resolution	Hourly historical data
Coverage	Global, 1940 to present
Usage	Fetched per job using location coordinates from the dataset's region

5.2 Holiday/Calendar — Nager.Date

Attribute	Detail
Provider	Nager.Date (date.nager.at)
Cost	Free, open source
Coverage	110+ countries, public holidays
Usage	Fetched once per country/year combination, cached in DB
Extension	Custom shutdown/vacation dates can be added post-MVP

6. Frontend MVP

The MVP frontend is intentionally minimal. Its only job is to let a user upload a file and view/download the results. No complex UI, no dashboards, no multi-user management.

6.1 User Flow

- Land on page → see drag-and-drop upload zone
- Drop or select file (CSV, Excel, MSCONS)
- Optional: provide timezone hint and dataset label
- Submit → see processing status with stage indicator
- On completion → see summary card (peak load, annual consumption, quality score)
- Download forecast as CSV or Excel

6.2 Frontend Pages / Views

View	Description
Upload	Drag & drop zone, optional fields, submit button
Processing	Job status with live polling, stage progress indicator
Results	Summary metrics, quality report overview, download buttons
Error	Clear error message with option to re-upload

7. Open Questions for Client Review

The following items require client input before or during development. Please review and confirm or provide guidance on each.

#	Question	Why It Matters
1	Which primary markets/countries should be supported at launch? (e.g. DE, GB, FR)	Determines default timezone, holiday calendar, and MSCONS variant
2	Is a user account/login required in the MVP, or is API-key-only access sufficient?	Affects auth implementation scope significantly
3	Should the 8,760 forecast vector be based on a specific future year, or always the next calendar year?	Affects Prophet model configuration
4	Is the weather enrichment mandatory for processing, or optional enhancement?	Affects what happens if weather API is unavailable
5	What is the expected maximum file size for uploads? (impacts GCS config and timeout settings)	Infrastructure sizing
6	Should imputed values be visible flagged in the download output?	Affects output schema
7	Is there an existing product name or brand for the repo and API naming?	Affects naming of repo, endpoints, docs
8	Are there specific confidence interval widths required? (default: 95%)	Prophet CI configuration

End of Document — Prepared for Client Review

Please confirm, amend, or raise questions against any section above.